



American International University-Bangladesh (AIUB)

**Department of Computer Science
Faculty of Science & Technology (FST)**

Course name: Introduction to Data Science

Semester: Summer 22-23

Section: A

Finalterm Project

Submitted To: Tohedul Islam

Assistant Professor , FACULTY OF SCIENCE & TECHNOLOGY
AMERICAN INTERNATIONAL UNIVERSITY–BANGLADESH

Submitted By:

Student Name	ID
MD. KAMRUZZAMAN	20-44222-3

Date Of Submission: August 16, 2023

About the dataset: This dataset is originally from the National Institute of Diabetes and Digestive and Kidney Diseases. The objective of the dataset is to diagnostically predict whether a patient has diabetes, based on certain diagnostic measurements included in the dataset. Several constraints were placed on the selection of these instances from a larger database. All patients here are females.

From the data set in the (.csv) file we can find several variables, some of them are independent (several medical predictor variables) and only one target dependent variable (Outcome).

Here are the details of dataset attributes:

Pregnancies - To express the Number of pregnancies.

Glucose - To express the Glucose level in blood .

BloodPressure - To express the Blood pressure measurement.

SkinThickness - To express the thickness of the skin.

Insulin - To express the Insulin level in blood.

BMI - To express the Body mass index.

DiabetesPedigreeFunction - To express the Diabetes percentage.

Age - To express the age.

Outcome - To express the final result 1 is Yes and 0 is No.

Link of the Dataset: <https://www.kaggle.com/datasets/whenamancodes/predict-diabilities>

Data Import and view

Code:

```
mydata<-read.csv("E:/Academic/AIUB/Semester 9/Data  
Science/Final/Project/MyProject/Dataset/diabetes.csv",header=TRUE,sep=",")  
mydata
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
1	6	148	72	35	0	33.6	0.627	50	1
2	1	85	66	29	0	26.6	0.351	31	0
3	8	183	64	0	0	23.3	0.672	32	1
4	1	89	66	23	94	28.1	0.167	21	0
5	0	137	40	35	168	43.1	2.288	33	1
6	5	116	74	0	0	25.6	0.201	30	0
7	3	78	50	32	88	31.0	0.248	26	1
8	10	115	0	0	0	35.3	0.134	29	0
9	2	197	70	45	543	30.5	0.158	53	1
10	8	125	96	0	0	0.0	0.232	54	1
11	4	110	92	0	0	37.6	0.191	30	0
12	10	168	74	0	0	38.0	0.537	34	1
13	10	139	80	0	0	27.1	1.441	57	0
14	1	189	60	23	846	30.1	0.398	59	1
15	5	166	72	19	175	25.8	0.587	51	1
16	7	100	0	0	0	30.0	0.484	32	1
17	0	118	84	47	230	45.8	0.551	31	1
18	7	107	74	0	0	29.6	0.254	31	1
19	1	103	30	38	83	43.3	0.183	33	0
20	1	115	70	30	96	34.6	0.529	32	1
21	3	126	88	41	235	39.3	0.704	27	0
22	8	99	84	0	0	35.4	0.388	50	0
23	7	196	90	0	0	39.8	0.451	41	1
24	9	119	80	35	0	29.0	0.263	29	1
25	11	143	94	33	146	36.6	0.254	51	1

Description: Data are read from a CSV file and store it in a variable named mydata.

Summary of dataset

Code:

```
mydata<-read.csv("E:/Academic/AIUB/Semester 9/Data
Science/Final/Project/MyProject/Dataset/diabetes.csv",header=TRUE,sep=",")
mydata
summary(mydata)
```

```
> summary(mydata)
      Pregnancies      Glucose      BloodPressure      SkinThickness      Insulin      BMI      DiabetesPedigreeFunction      Age
Min.   : 0.000   Min.   : 0.0   Min.   : 0.00   Min.   : 0.00   Min.   : 0.0   Min.   : 0.00   Min.   :0.0780   Min.   :21.00
1st Qu.: 1.000   1st Qu.: 99.0   1st Qu.: 62.00   1st Qu.: 0.00   1st Qu.: 0.0   1st Qu.:27.30   1st Qu.:0.2437   1st Qu.:24.00
Median : 3.000   Median :117.0   Median : 72.00   Median :23.00   Median :30.5   Median :32.00   Median :0.3725   Median :29.00
Mean   : 3.845   Mean   :120.9   Mean   : 69.11   Mean   :20.54   Mean   : 79.8   Mean   :31.99   Mean   :0.4719   Mean   :33.24
3rd Qu.: 6.000   3rd Qu.:140.2   3rd Qu.: 80.00   3rd Qu.:32.00   3rd Qu.:127.2   3rd Qu.:36.60   3rd Qu.:0.6262   3rd Qu.:41.00
Max.   :17.000   Max.   :199.0   Max.   :122.00   Max.   :99.00   Max.   :846.0   Max.   :67.10   Max.   :2.4200   Max.   :81.00
      Outcome
Min.   :0.000
1st Qu.:0.000
Median :0.000
Mean   :0.349
3rd Qu.:1.000
Max.   :1.000
```

Description: This generates a summary of the mydata data using the summary() function. It's a useful function for initial data exploration and understanding the characteristics of your dataset.

Display dataset

Code: mydata<-read.csv("E:/Academic/AIUB/Semester 9/Data

```
Science/Final/Project/MyProject/Dataset/diabetes.csv",header=TRUE,sep=",")
```

```
mydata
```

```
str(mydata)
```

```
> str(mydata)
'data.frame':   768 obs. of  9 variables:
 $ Pregnancies      : int  6 1 8 1 0 5 3 10 2 8 ...
 $ Glucose          : int  148 85 183 89 137 116 78 115 197 125 ...
 $ BloodPressure    : int  72 66 64 66 40 74 50 0 70 96 ...
 $ SkinThickness    : int  35 29 0 23 35 0 32 0 45 0 ...
 $ Insulin          : int  0 0 0 94 168 0 88 0 543 0 ...
 $ BMI              : num  33.6 26.6 23.3 28.1 43.1 25.6 31 35.3 30.5 0 ...
 $ DiabetesPedigreeFunction: num  0.627 0.351 0.672 0.167 2.288 ...
 $ Age              : int  50 31 32 21 33 30 26 29 53 54 ...
 $ Outcome          : int  1 0 1 0 1 0 1 0 1 1 ...
```

Description: “str” displays structures of R objects. “str” used for displaying the contents of the data. str () is an alternative function to display the summary of the output produced, especially when the data set is huge.

Missing value detection

Code: `is.na(mydata)`

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
[1,]	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE
[2,]	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE
[3,]	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE
[4,]	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE
[5,]	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE
[6,]	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE
[7,]	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE
[8,]	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE
[9,]	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE
[10,]	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE
[11,]	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE
[12,]	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE
[13,]	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE
[14,]	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE
[15,]	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE
[16,]	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE
[17,]	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE
[18,]	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE
[19,]	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE
[20,]	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE
[21,]	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE
[22,]	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE

Description: The `is.na()` function is used to identify missing values in a data frame or vector. It returns a logical vector of the same length as the input data, where each element is TRUE if the corresponding element in the data frame or vector is missing, and FALSE otherwise.

Missing Value count in each column

Code: `colSums(is.na(mydata))`

```
> colSums(is.na(mydata))
Pregnancies      0      Glucose      0      BloodPressure      0      SkinThickness      0      Insulin      0
BMI      0      DiabetesPedigreeFunction      0      Age      0      Outcome      0
```

Description: The `colSums()` function is used to calculate the sum of elements for each column.

Normalization

Code: `mydata`

```
normalize <- function(x) {
  return ((x - min(x)) / (max(x) - min(x)))
}
mydata <- as.data.frame(lapply(mydata[1:9], normalize))
mydata
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
1	0.35294118	0.7437186	0.5901639	0.3535354	0.00000000	0.5007452	0.23441503	0.48333333	1
2	0.05882353	0.4271357	0.5409836	0.2929293	0.00000000	0.3964232	0.11656704	0.16666667	0
3	0.47058824	0.9195980	0.5245902	0.00000000	0.00000000	0.3472429	0.25362938	0.18333333	1
4	0.05882353	0.4472362	0.5409836	0.2323232	0.11111111	0.4187779	0.03800171	0.00000000	0
5	0.00000000	0.6884422	0.3278689	0.3535354	0.19858156	0.6423249	0.94363792	0.20000000	1
6	0.29411765	0.5829146	0.6065574	0.00000000	0.00000000	0.3815201	0.05251921	0.15000000	0
7	0.17647059	0.3919598	0.4098361	0.3232323	0.10401891	0.4619970	0.07258753	0.08333333	1
8	0.58823529	0.5778894	0.00000000	0.00000000	0.00000000	0.5260805	0.02391119	0.13333333	0
9	0.11764706	0.9899497	0.5737705	0.4545455	0.64184397	0.4545455	0.03415884	0.53333333	1
10	0.47058824	0.6281407	0.7868852	0.00000000	0.00000000	0.00000000	0.06575576	0.55000000	1
11	0.23529412	0.5527638	0.7540984	0.00000000	0.00000000	0.5603577	0.04824936	0.15000000	0
12	0.58823529	0.8442311	0.5065574	0.00000000	0.00000000	0.5663188	0.10508634	0.31666667	1

Description: My implemented normalization method known as "min-max normalization". This type of normalization scales the values of each variable to a range between 0 and 1 based on their minimum and maximum values. Data normalization can be performed using various methods, including manual calculations or through libraries like caret.

Correlation

Code: correlation<- cor(mydata)
correlation

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
Pregnancies	1.00000000	0.12945867	0.14128198	-0.08167177	-0.07353461	0.01768309	-0.03352267	0.54434123	
Glucose	0.12945867	1.00000000	0.15258959	0.05732789	0.33135711	0.22107107	0.13733730	0.26351432	
BloodPressure	0.14128198	0.15258959	1.00000000	0.20737054	0.08893338	0.28180529	0.04126495	0.23952795	
SkinThickness	-0.08167177	0.05732789	0.20737054	1.00000000	0.43678257	0.39257320	0.18392757	-0.11397026	
Insulin	-0.07353461	0.33135711	0.08893338	0.43678257	1.00000000	0.19785906	0.18507093	-0.04216295	
BMI	0.01768309	0.22107107	0.28180529	0.39257320	0.19785906	1.00000000	0.14064695	0.03624187	
DiabetesPedigreeFunction	-0.03352267	0.13733730	0.04126495	0.18392757	0.18507093	0.14064695	1.00000000	0.03356131	
Age	0.54434123	0.26351432	0.23952795	-0.11397026	-0.04216295	0.03624187	0.03356131	1.00000000	
Outcome	0.22189815	0.46658140	0.06506836	0.07475223	0.13054795	0.29269466	0.17384407	0.23835598	
Outcome									
Pregnancies									
Glucose									
BloodPressure									
SkinThickness									
Insulin									
BMI									
DiabetesPedigreeFunction									
Age									
Outcome									

Description: Calculated the correlation between variables using the cor() function. This function computes the correlation coefficient between pairs of variables in a data frame.

Correlation Plotting

Code: install.packages("corrplot")
library(corrplot)
cor_matrix <- cor(mydata)
corrplot(cor_matrix, method = "number")



Description: Correlation plotting for better visualization.

Dividing dataset into training and test set

Code: `train_indices<-sample(nrow(mydata), 0.8 * nrow(mydata))`

`train_indices`

`train_data<-mydata[train_indices,]`

`train_data`

`test_data<-mydata[-train_indices,]`

`test_data`

```
> train_data
  Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin    BMI  DiabetesPedigreeFunction    Age  Outcome
435  0.05882353 0.4522613    0.5573770    0.08080808 0.00000000 0.3651267    0.452604611 0.25000000    0
19   0.05882353 0.5175879    0.2459016    0.38383838 0.09810875 0.6453055    0.044833476 0.20000000    0
410  0.05882353 0.8643216    0.5573770    0.49494949 0.68439716 0.6318927    0.266438941 0.11666667    1
102  0.05882353 0.7587940    0.4918033    0.00000000 0.00000000 0.3889717    0.043125534 0.01666667    0
504  0.41176471 0.4723618    0.5245902    0.25252525 0.09338061 0.4962742    0.281810418 0.33333333    0
420  0.17647059 0.6482412    0.5245902    0.29292929 0.13593381 0.3934426    0.060204953 0.11666667    1
364  0.23529412 0.7336683    0.6393443    0.00000000 0.00000000 0.5737705    0.188727583 0.76666667    1
341  0.05882353 0.6532663    0.5737705    0.13131313 0.12411348 0.3859911    0.168232280 0.01666667    0
114  0.23529412 0.3819095    0.5081967    0.00000000 0.00000000 0.5067064    0.133646456 0.06666667    0
654  0.11764706 0.6030151    0.4426230    0.00000000 0.00000000 0.3994039    0.160973527 0.10000000    0
```

```
> test_data
  Pregnancies  Glucose BloodPressure SkinThickness  Insulin      BMI DiabetesPedigreeFunction   Age Outcome
13  0.58823529 0.6984925    0.6557377    0.00000000 0.00000000 0.4038748    0.58198121 0.60000000    0
14  0.05882353 0.9497487    0.4918033    0.23232323 1.00000000 0.4485842    0.13663535 0.63333333    1
26  0.58823529 0.6281407    0.5737705    0.26262626 0.13593381 0.4634873    0.05422716 0.33333333    1
28  0.05882353 0.4874372    0.5409836    0.15151515 0.16548463 0.3457526    0.17463706 0.01666667    0
45  0.41176471 0.7989950    0.5245902    0.00000000 0.00000000 0.4083458    0.09222886 0.31666667    0
48  0.11764706 0.3567839    0.5737705    0.27272727 0.00000000 0.4172876    0.21690863 0.01666667    0
49  0.41176471 0.5175879    0.5409836    0.32323232 0.00000000 0.5827124    0.11357814 0.16666667    1
52  0.05882353 0.5075377    0.4098361    0.15151515 0.04255319 0.3606557    0.19128950 0.08333333    0
53  0.29411765 0.4422111    0.5409836    0.21212121 0.02718676 0.3636364    0.11272417 0.15000000    0
--  - - - - -
```

Description: To split our data into 80% for training and 20% for testing set, sample() function used.

Find the total number of observation and value of K

Code: NROW(train_data)
sqrt(NROW(train_data))

```
> NROW(train_data)
[1] 614
> sqrt(NROW(train_data))
[1] 24.77902
```

Description: Calculated the number of observations in the training data set. The reason for doing this is to initialize the value of 'K' in the KNN model. One of the ways to find the optimal K value is to calculate the square root of the total number of observations in the data set. This square root will give the 'K' value.

Accuracy of dividing the data into training and test set approach

Code: knn.24 <- knn(train = train_data[, -9], test = test_data[, -9], cl = train_data\$Outcome, k = 24)
knn.25 <- knn(train = train_data[, -9], test = test_data[, -9], cl = train_data\$Outcome, k = 25)
Accuracy.24 <- 100 * sum(test_data\$Outcome == knn.24) / nrow(test_data)
Accuracy.24
Accuracy.25 <- 100 * sum(test_data\$Outcome == knn.25) / nrow(test_data)
Accuracy.25
> Accuracy.24
[1] 74.67532
> Accuracy.25 <- 100 * sum(test_data\$Outcome == knn.25) / nrow(test_data)
> Accuracy.25
[1] 75.32468

Description: Calculated the accuracy of dividing the data into training and test set approach.

Confusion matrix of dividing the data into training and test set approach

Code: For K=24
confusion_matrix.24 <- table(knn.24, test_data\$Outcome)
confusion_matrix.24
confusion_matrix_summary.24 <- confusionMatrix(confusion_matrix.24)

```

confusion_matrix_summary.24
recall.24 <- confusion_matrix_summary.24$byClass["Sensitivity"]
recall.24
precision.24 <- confusion_matrix_summary.24$byClass["Pos Pred Value"]
precision.24

```

```

knn.24  0  1
      0 85 21
      1 18 30

      Accuracy : 0.7468
      95% CI : (0.6705, 0.8133)
      No Information Rate : 0.6688
      P-Value [Acc > NIR] : 0.02264

      Kappa : 0.4197

      McNemar's Test P-Value : 0.74877

      Sensitivity : 0.8252
      Specificity : 0.5882
      Pos Pred Value : 0.8019
      Neg Pred Value : 0.6250
      Prevalence : 0.6688
      Detection Rate : 0.5519
      Detection Prevalence : 0.6883
      Balanced Accuracy : 0.7067

      'Positive' Class : 0
      > recall.24
      Sensitivity
      0.8252427
      > precision.24
      Pos Pred Value
      0.8018868

```

Code: For K=25

```

confusion_matrix.25 <- table(knn.25, test_data$Outcome)
confusion_matrix.25
confusion_matrix_summary.25 <- confusionMatrix(confusion_matrix.25)
confusion_matrix_summary.25

```



```

knn.25  0  1
       0 86 21
       1 17 30

Accuracy : 0.7532
 95% CI : (0.6774, 0.8191)
No Information Rate : 0.6688
P-Value [Acc > NIR] : 0.01456

Kappa : 0.4317

McNemar's Test P-Value : 0.62650

Sensitivity : 0.8350
Specificity : 0.5882
Pos Pred Value : 0.8037
Neg Pred Value : 0.6383
Prevalence : 0.6688
Detection Rate : 0.5584
Detection Prevalence : 0.6948
Balanced Accuracy : 0.7116

'Positive' Class : 0
> recall.25
Sensitivity
0.8349515
> precision.25
Pos Pred Value
0.8037383

```

Description: Confusion matrix generated to calculate the accuracy of the KNN model with K value set to 24 and 25 respectively. Recall and Precision value also calculated for this KNN classifier.

10-fold cross validation

Code: `ctrl <- trainControl(method = "cv", number = 10)`
`knn_model <- train(`
 `Outcome ~ .,`
 `data = mydata,`
 `method = "knn",`
 `trControl = ctrl,`
 `preProcess = c("center", "scale"),`
 `tuneGrid = data.frame(k = 10)`
`)`
`knn_model`

```

> knn_model
k-Nearest Neighbors

768 samples
  8 predictor

Pre-processing: centered (8), scaled (8)
Resampling: Cross-Validated (10 fold)
Summary of sample sizes: 691, 692, 691, 691, 691, 691, ...
Resampling results:

      RMSE      Rsquared    MAE
0.4193702  0.2429772  0.3264374

Tuning parameter 'k' was held constant at a value of 10

```

Description: 10-fold cross validation approach implemented here and Root Mean Square Error, R-squared (Coefficient of Determination), Mean Absolute Error calculated.

10-fold cross validation Confusion matrix

Code: ctrl <- trainControl(method = "cv", number = 10)

```

knn_model <- train(
  Outcome ~ .,
  data = mydata,
  method = "knn",
  trControl = ctrl,
  preProcess = c("center", "scale"),
  tuneGrid = data.frame(k = 10)
)
knn_model

conf_matrix

```

```

> conf_matrix
Confusion Matrix and Statistics

      Reference
Prediction 0  1
0  437 108
1   63 160

      Accuracy : 0.7773
      95% CI   : (0.7462, 0.8063)
No Information Rate : 0.651
P-Value [Acc > NIR] : 1.833e-14

      Kappa : 0.4901
McNemar's Test P-Value : 0.0007661

      Sensitivity : 0.8740
      Specificity : 0.5970
      Pos Pred Value : 0.8018
      Neg Pred Value : 0.7175
      Prevalence : 0.6510
      Detection Rate : 0.5690
      Detection Prevalence : 0.7096
      Balanced Accuracy : 0.7355

      'Positive' Class : 0
> recall <- conf_matrix$byClass["Sensitivity"]
> recall
Sensitivity
0.874
> precision <- conf_matrix$byClass["Precision"]
> precision
Precision
0.8018349

```

Description: Confusion matrix generated to calculate the accuracy of the KNN model with K value set to 10. Recall and Precision value also calculated for this KNN classifier.